

Daily Trip Duration Analysis of Uber Rides in NYC Using Time Series Methods

Farooq Mahmud

August 7, 2025

Contents

1	Introduction	2
2	Preliminary Analysis	2
2.1	Data Preparation and Description	2
2.2	Initial Visualization and Trends	3
2.3	Annotated Time Series and Weekly Seasonality	4
2.4	Stationarity Assessment	4
2.5	Autocorrelation and Partial Autocorrelation Analysis	4
2.6	Reconciling ADF and ACF/PACF Findings	5
3	Secondary Analysis and Modeling	5
3.1	Model Selection and Diagnostics	5
3.2	Model Specification	6
3.3	Model Estimation and Comparison	6
3.4	Diagnostic Checking	6
3.5	Model Selection Summary	9
3.6	Fitted Model Equation	9
3.7	Interpretation of the ARIMA Model	9
3.8	Forecasting	10
4	Discussion and Conclusion	11
A	Download Instructions	13
B	R Time Series Modeling and Visualization Code	13
C	PySpark Preprocessing Code	15

1 Introduction

Urban mobility has undergone a transformative shift with the proliferation of ride-hailing services such as Uber. In dense metropolitan areas such as New York City, these services generate massive amounts of data that can reveal important patterns in travel behavior. This project proposes a time series analysis of average daily Uber trip durations in New York City throughout the year 2024, using publicly available high-volume for-hire vehicle (FHV) trip data.

The dataset is sourced from the NYC Taxi and Limousine Commission (TLC) and contains more than 200 million records for 2024 alone. For the purpose of this analysis, a subset of the data has been preprocessed using PySpark to compute the average duration of Uber trips per day, resulting in 365 data points. Time series modeling and visualization were performed exclusively in R. The dataset is accessible through the NYC TLC portal [2].

The main objective of this project is to explore the temporal behavior of Uber trip durations and apply time series modeling to understand and forecast future values. Specific goals include examining stationarity, identifying appropriate models through diagnostics, and evaluating the accuracy of forecasts. A review of previous research indicates that ride-hailing data have been used extensively for congestion, demand prediction, and fleet optimization, but less so for temporal duration modeling at a daily resolution.

Two references guiding this work are (1) Tandon et al. (2024), which evaluates the effect of congestion pricing on ride-hailing behavior in New York City [3] and (2) Ma et al. (2024), which explores congestion-aware scheduling of electric ride-hailing fleets [1].

This study contributes by applying time series methods to a large-scale, real-world dataset to understand trends and variability in urban travel times.

All R code used in the analysis is provided in Appendix B. References to code listings (e.g., Listing 1) throughout this paper correspond to entries in the Appendix.

2 Preliminary Analysis

2.1 Data Preparation and Description

The dataset used in this study consists of the average daily trip duration (in minutes) for Uber trips in New York City throughout 2024. This time series was constructed by aggregating more than 200 million ride records published by the New York Taxi and Limousine Commission. A PySpark pipeline was used to compute the average duration per day, resulting in a time series of 366 observations (one for each day in the leap year 2024). The complete code is included in Appendix C.

The decision to use the average daily trip duration as the primary variable was motivated by both practical and analytical considerations. From a practical point of view, the original dataset comprises more than 200 million individual




Figure 1: Daily average Uber trip duration in NYC for 2024.

trip records, which would pose significant computational challenges for time series modeling and visualization. Aggregating to daily averages drastically reduces the data size, while preserving temporal patterns of interest.

Analytically, daily aggregation smooths out short-term noise (e.g., from trip-level variability) and highlights broader trends, such as weekly or seasonal cycles. This level of granularity strikes a balance between detail and explanatory power, making it suitable for time-series methods such as ARIMA. It also allows for meaningful comparisons between days, making the results more actionable to understand congestion and mobility behavior.

2.2 Initial Visualization and Trends

To begin the exploratory phase, a time series line plot of the daily average trip durations was generated (Listing 1). Figure 1 displays this unannotated time series. Several key patterns emerge, even without additional labeling. The plot exhibits moderate variability throughout the year with visible periodic fluctuations. A general upward trend is observed during the summer months, suggesting seasonal effects. There also appear to be isolated spikes and drops, which may correspond to holidays, weather events, or changes in urban mobility behavior.




Figure 2: Daily average Uber trip duration in NYC for 2024, with weekends and holidays highlighted.

2.3 Annotated Time Series and Weekly Seasonality

In an attempt to expose more patterns, a second version of the time series plot was created with Fridays, Saturdays, Sundays, and the major holidays in the United States highlighted (Figure 2, Listing 2). This annotated visualization reveals more structured weekly seasonality, with higher average durations occurring more frequently on weekends. Holidays such as Memorial Day, Independence Day, and Christmas are associated with pronounced decreases in duration.

2.4 Stationarity Assessment

To determine whether the daily average Uber trip duration series is stationary, I performed the Augmented Dickey-Fuller (ADF) test (Listing 3). The ADF test yielded a p -value of 0.04612. Because the p -value is less than 0.05, I reject the null hypothesis at the significance level $\alpha=0.05$. Therefore, there is statistical evidence to suggest that the time series is stationary.

2.5 Autocorrelation and Partial Autocorrelation Analysis

The ACF plot (Figure 3, Listing 4) shows a slow and gradual decline, consistent with a non-stationary series. This pattern suggests the presence of a trend and the need to differentiate the series to induce stationarity. Significant autocorrelation at multiple lags may also indicate latent seasonality.




Figure 3: Autocorrelation Function (ACF) plot of daily average trip durations. The plot shows a slow decay, suggesting non-stationarity and potential seasonality in the data.

The PACF plot (Figure 4, Listing 5) shows a large spike at lag 1, followed by a tapering tail of smaller spikes. This pattern suggests the existence of an autoregressive (AR) structure in the series. The combination of these findings supports the consideration of ARIMA modeling, potentially with seasonal differencing.

2.6 Reconciling ADF and ACF/PACF Findings

Although the ADF test returned a p -value of 0.04612, suggesting stationarity at significance level $\alpha=0.05$, the ACF and PACF plots exhibit patterns indicative of non-stationarity, including a slow decay in autocorrelation. This discrepancy implies that the series may be only marginally stationary or trend-stationary. To address this, further differencing is warranted in the model's subsequent development.

3 Secondary Analysis and Modeling

3.1 Model Selection and Diagnostics

To identify the most appropriate model for forecasting the average duration of Uber trips, I conducted a comprehensive model selection and diagnostic process.




Figure 4: Partial Autocorrelation Function (PACF) plot of daily average trip durations. The sharp drop after the first few lags indicates the presence of autoregressive components.

3.2 Model Specification

I first applied first-order differencing to the original time series to achieve stationarity, confirmed visually and through the ADF test. The ACF PACF plots were then examined to guide the selection of autoregressive (AR) and moving average (MA) terms. Based on these tools and the output of `auto.arima`, I selected the following candidate models for further analysis:

- ARIMA(1,1,1)
- ARIMA(2,1,1)
- ARIMA(3,1,2)

3.3 Model Estimation and Comparison

Each model was fitted and their performance were evaluated using the Akaike Information Criterion (AIC) and Bayesian Information Criterion (BIC) (Listing 7). The results are summarized in Table 1.

ARIMA(3,1,2) had the lowest AIC and BIC values, suggesting that it is the most parsimonious and best-fitting model among the three.

3.4 Diagnostic Checking

I conducted residual diagnostics for each model, including residual plots (Figures 5, 6, and 7 and the Ljung-Box test for autocorrelation (Listing 8). Table 2

Table 1: Model comparison using AIC and BIC

Model	AIC	BIC
ARIMA(1,1,1)	1028.02	1039.72
ARIMA(2,1,1)	1019.60	1035.20
ARIMA(3,1,2)	962.08	985.48




Figure 5: Residual diagnostics for ARIMA(1,1,1): time plot (top), ACF of residuals (bottom left), and histogram with normal distribution curve (bottom right).

presents the results of the Ljung box tests.

Table 2: Ljung-Box test results on model residuals

Model	Q* Statistic	df	p-value
ARIMA(1,1,1)	810.27	71	$< 2.2 \times 10^{-16}$
ARIMA(2,1,1)	477.14	70	$< 2.2 \times 10^{-16}$
ARIMA(3,1,2)	88.01	68	0.05184

The ARIMA(3,1,2) model was the only one to pass the Ljung-Box test at significance level $\alpha = 0.05$, with a p -value of 0.05184, indicating no significant residual autocorrelation. The residuals also appeared to be normally distributed. ACF plots of the residuals confirmed the absence of strong autocorrelation.

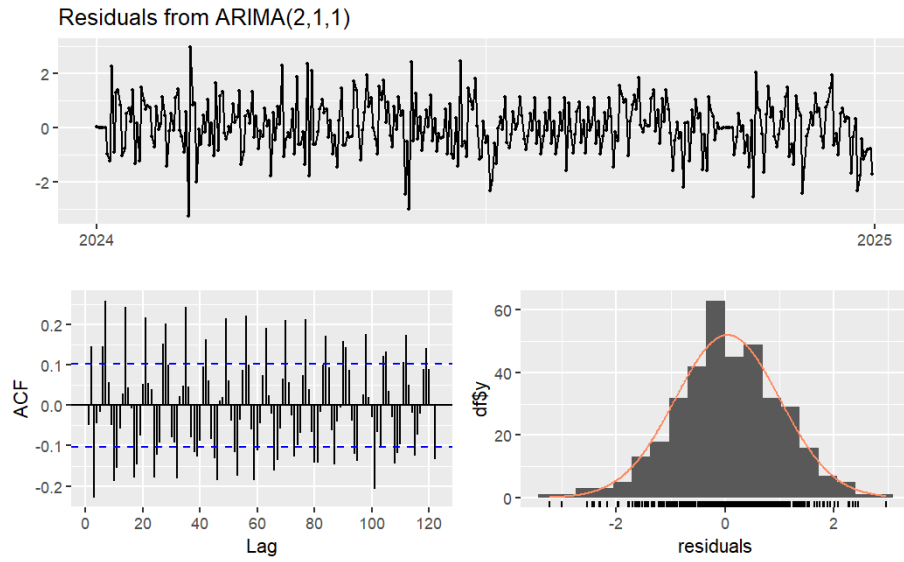


Figure 6: Residual diagnostics for ARIMA(2,1,1): time plot (top), ACF of residuals (bottom left), and histogram with normal distribution curve (bottom right).




Figure 7: Residual diagnostics for ARIMA(3,1,2): time plot (top), ACF of residuals (bottom left), and histogram with normal distribution curve (bottom right).

3.5 Model Selection Summary

Based on information criterion, residual diagnostics, and parsimony, I selected ARIMA(3,1,2) as the final model for forecasting. It captures the temporal dependencies in the data and produces uncorrelated, approximately normal residuals, making it suitable for short-term prediction.

ARIMA (3,1,2)

Coefficients:

ar1	ar2	ar3	ma1	ma2
0.8166	-0.4500	-0.3559	-1.2271	0.6940

Figure 8: Estimated ARIMA(3,1,2) model coefficients.

The model includes three autoregressive terms (AR) and two moving average terms (MA), with one order of differencing. The residuals from this model were analyzed using diagnostic plots and the Ljung-Box test to assess the adequacy of the model.

3.6 Fitted Model Equation

The fitted ARIMA(3,1,2) model can be written as shown in Equation 1.

$$Y_t = Y_{t-1} + 0.8166(Y_{t-1} - Y_{t-2}) - 0.4500(Y_{t-2} - Y_{t-3}) - 0.3559(Y_{t-3} - Y_{t-4}) + \varepsilon_t - 1.2271\varepsilon_{t-1} + 0.6940\varepsilon_{t-2} \quad (1)$$

where Y_t denotes the average trip duration at time t and ε_t is a white noise error term.

3.7 Interpretation of the ARIMA Model

The fitted ARIMA(3,1,2) model suggests that the average Uber trip duration on a given day is influenced by a combination of past values and past random fluctuations.

The **autoregressive (AR)** component means that the duration is partly predicted using the trip durations of the last three days. If durations have been increasing or decreasing in recent days, this trend is likely to continue in the short term.

The **integrated (I)** component—represented by the differencing step - accounts for nonstationarity, helping the model focus on changes from one day to the next rather than absolute levels.

The **moving average (MA)** component adjusts the predictions based on the last two prediction errors. For example, if a previous day's duration was

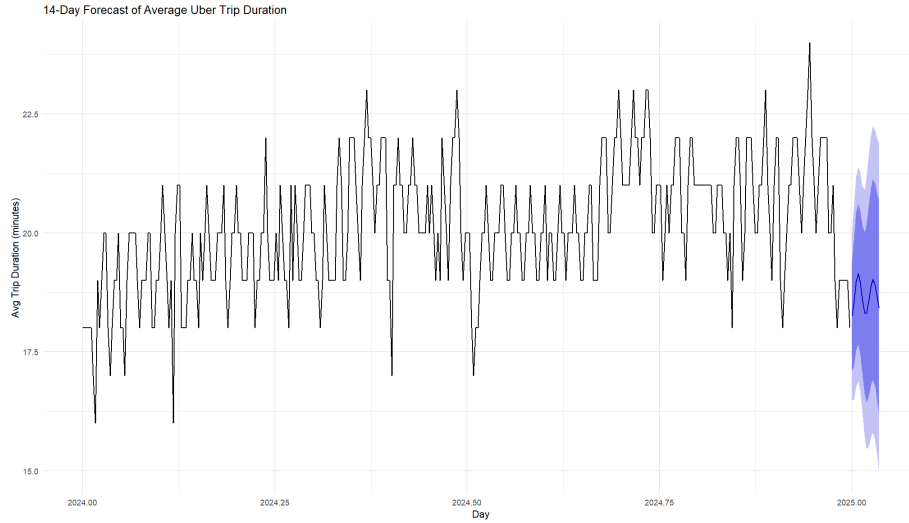


Figure 9: 14-day forecast of average Uber trip duration with 80% and 95% confidence intervals.

unusually high or low due to unexpected traffic or a weather event, the MA terms help correct for that anomaly in future predictions.

Together, these components allow the ARIMA model to capture both consistent temporal patterns and unexpected variations in ride durations, making it a practical tool to forecast and understand urban mobility trends over time.

3.8 Forecasting

To evaluate the predictive utility of the fitted model, I generated a 14-day forecast using the ARIMA(3,1,2) model (Listing 9). Figure 9 shows the projected average trip duration along with the 80% and 95% confidence intervals (Listing 10).

The forecast suggests that the average travel duration is expected to stabilize in the range of 18 to 19 minutes. Although point forecasts are relatively flat, the widening confidence intervals reflect increased uncertainty over time. The forecasted values for the next 14 days, along with their 80% and 95% prediction intervals (Listing 12), are shown in Table 3.

Residual diagnostics (Figure 10, Listing 11) indicate that the residuals are approximately normally distributed and do not exhibit significant autocorrelation, as confirmed by the Ljung-Box test ($p = 0.0518$), suggesting that the model adequately captures the structure in the data.

Table 3: 14-day forecast of average trip durations with 80% and 95% prediction intervals.

Date	Forecast	Lo 80	Hi 80	Lo 95	Hi 95
2025-01-01	18.25	17.11	19.40	16.50	20.00
2025-01-02	18.52	17.19	19.85	16.48	20.55
2025-01-03	18.98	17.53	20.42	16.76	21.19
2025-01-04	19.14	17.67	20.62	16.89	21.40
2025-01-05	18.98	17.48	20.48	16.68	21.27
2025-01-06	18.60	17.04	20.16	16.22	20.99
2025-01-07	18.31	16.62	20.00	15.73	20.90
2025-01-08	18.30	16.44	20.17	15.46	21.15
2025-01-09	18.56	16.56	20.56	15.50	21.62
2025-01-10	18.88	16.80	20.95	15.71	22.04
2025-01-11	19.02	16.92	21.13	15.80	22.24
2025-01-12	18.91	16.78	21.04	15.65	22.17
2025-01-13	18.64	16.46	20.81	15.31	21.97
2025-01-14	18.42	16.15	20.68	14.95	21.88

4 Discussion and Conclusion

This project set out to analyze temporal trends in Uber trip durations using New York City’s High-Volume For-Hire Vehicle (FHV) dataset for 2024. The analysis focused on daily average trip durations and applied time series modeling techniques to identify patterns and forecast future behavior.

The ARIMA(3,1,2) model fitted to the data demonstrated a reasonably good fit, with low residual autocorrelation and a Ljung-Box test p -value close to significance level $\alpha = 0.05$. The model successfully captured temporal dynamics and was informed by stationarity checks, ACF, PACF, AIC, and BIC analysis. The results indicate that the duration of Uber rides exhibits regular fluctuations, with weekends and most holidays showing elevated travel times, likely due to increased congestion or changes in demand patterns.

The fitted ARIMA model was also used to generate a 14-day forecast of average trip durations. The forecast suggests continued fluctuations with point predictions that ranged from approximately 18 to 19 minutes. The 95% prediction intervals provide a reasonable range of uncertainty, with the lower bounds staying above 16 minutes and the upper bounds exceeding 21 minutes at times. This demonstrates the model’s utility for short-term forecasting of ride duration trends.

Importantly, the structure of the ARIMA model provides interpretable insights. The autoregressive (AR) terms reflect the persistence of duration trends over three days, while the moving average (MA) terms capture the impact of recent anomalies such as sudden spikes or drops in trip duration due to unpredictable events such as traffic disruptions or weather. The presence of two MA terms allows the model to learn from forecasting errors over the past two days,




Figure 10: Residual diagnostics for the 14-day ARIMA forecast: time plot (top), ACF of residuals (bottom left), and histogram with normal distribution curve (bottom right).

making the forecasts more adaptive and robust to short-term volatility.

One strength of the analysis was the use of big data tools (PySpark) to preprocess more than 200 million records into a manageable time series format. This preprocessing enabled meaningful aggregation and reduced the computational burden on R, which was used for modeling and visualization. Visualizing holidays and weekends also helped contextualize anomalies and recurring patterns in the data.

However, the study has several limitations. The univariate ARIMA model focuses solely on past values and past errors in average trip duration, without explicitly incorporating external variables such as weather conditions, traffic incidents, or special events. Although moving average terms help absorb recent unexplained anomalies, they do not directly account for known external factors that could improve predictive accuracy. Additionally, aggregating data at the daily level smooths over intraday variations, potentially obscuring important temporal dynamics. Finally, the dataset includes only Uber trips, which may limit the generalization of the findings to the broader transportation or ride-hailing ecosystem.

Future work could include incorporating a multivariate approach to examine the relationships between trip duration, trip count, and distance. More granular models (for example, hourly data) or spatial analysis could also reveal deeper insights into congestion dynamics across neighborhoods.

In conclusion, this project demonstrates the feasibility and value of using large-scale ride-hailing data for time-series analysis. The findings support the

use of predictive models to anticipate congestion trends and could inform urban planning, fleet management, or dynamic pricing strategies.

A Download Instructions

The dataset, the R analysis script and the PySpark notebook are also available for download via a public GitHub release at <https://github.com/farooq-teqniqly/uwf-sta6856/releases/tag/final-project-v1.9/finalproject.zip>

Download and extract the ZIP file to access the CSV dataset and the code.

B R Time Series Modeling and Visualization Code

Listing 1: Loading and Plotting Time Series

```
1 library(ggplot2)
2
3 data <- read.csv(file.path(getwd(), "../data", "finalproject.csv"))
4 ts_data <- ts(data$avg_duration_min, frequency = 366, start = c(2024, 1))
5
6 library(ggplot2)
7 library(scales)
8
9 start_date <- as.Date("2024-01-01")
10 date_seq <- seq(from = start_date, by = "day", length.out = length(ts_data))
11
12 df <- data.frame(
13   Date = date_seq,
14   Duration = as.numeric(ts_data)
15 )
16
17 ggplot(df, aes(x = Date, y = Duration)) +
18   geom_line() +
19   labs(
20     title = "Daily Average Uber Trip Duration (2024)",
21     x = "Month",
22     y = "Average Duration (minutes)"
23   ) +
24   scale_x_date(
25     date_breaks = "1 month",
26     date_labels = "%b"
27   ) +
28   theme_minimal()
29
```

Listing 2: Plot Highlighting Weekends and Holidays

```
1 library(dplyr)
2
3 data$pickup_date <- as.Date(data$pickup_date)
4
5 data <- data %>%
6   mutate(weekday = weekdays(pickup_date))
7
8
9 fridays <- filter(data, weekday == "Friday")
10 saturdays <- filter(data, weekday == "Saturday")
11 sundays <- filter(data, weekday == "Sunday")
12
13 holidays <- data.frame(
14   date = as.Date(c(
15     "2024-01-01", # New Year's Day
16     "2024-01-15", # Martin Luther King Jr. Day
17     "2024-02-19", # Presidents' Day
18     "2024-05-27", # Memorial Day
19     "2024-07-04", # Independence Day
20     "2024-09-02", # Labor Day
21     "2024-10-14", # Columbus Day
22     "2024-11-11", # Veterans Day
23     "2024-11-28", # Thanksgiving
24   ))
25
```

```

24   "2024-12-25" # Christmas
25   )),
26   label = c(
27     "NewYear'sDay", "MLKJrDay", "Presidents'Day", "MemorialDay", "IndependenceDay",
28     "LaborDay", "ColumbusDay", "VeteransDay", "Thanksgiving", "Christmas"
29   )
30   )
31   )
32   holidays$avg_duration_min <- data$avg_duration_min[match(holidays$date, data$pickup_date)]
33
34   ggplot(data, aes(x = pickup_date, y = avg_duration_min)) +
35     geom_line(color = "gray", alpha = 0.5) +
36     geom_line(data = fridays, aes(color = "Friday")) +
37     geom_line(data = saturdays, aes(color = "Saturday")) +
38     geom_line(data = sundays, aes(color = "Sunday")) +
39     geom_point(data = holidays, aes(x = date, y = avg_duration_min),
40               color = "black", fill = "yellow", size = 2, shape = 21) +
41     geom_text(data = holidays, aes(x = date, y = avg_duration_min, label = label),
42              vjust = -1, size = 3, angle = 45) +
43     scale_color_manual(
44       values = c(
45         "Friday" = "#332288",
46         "Saturday" = "#EE7733",
47         "Sunday" = "#117733"
48       )
49     ) +
50     labs(
51       title = "DailyAverageTripDuration(2024)",
52       x = "Month", y = "AverageDuration(minutes)",
53       color = "Day"
54     ) +
55     theme_minimal() +
56     theme(
57       panel.grid.major = element_blank(),
58       panel.grid.minor = element_blank()
59     )

```

Listing 3: Augmented Dickey-Fuller (ADF) Test

```

1 library(tseries)
2 adf.test(ts_data)

```

Listing 4: ACF Plot

```

1 acf(ts_data, main = "ACFofDailyAverageDuration", lag.max = 150)

```

Listing 5: PACF Plot

```

1 pacf(ts_data, main = "PACFofDailyAverageDuration", lag.max = 150)

```

Listing 6: ARIMA Model

```

1 library(forecast)
2 diff_data <- diff(ts_data)
3 fit <- auto.arima(ts_data)
4 summary(fit)

```

Listing 7: Model Selection and Diagnostics

```

1 library(forecast)
2 diff_data <- diff(ts_data)
3
4 fit_1 <- Arima(ts_data, order = c(1,1,1))
5 fit_2 <- Arima(ts_data, order = c(2,1,1))
6 auto_fit <- auto.arima(ts_data)
7
8 model_comparison <- data.frame(
9   Model = c("ARIMA(1,1,1)", "ARIMA(2,1,1)", "ARIMA(3,1,2)"),
10   AIC = c(AIC(fit_1), AIC(fit_2), AIC(auto_fit)),
11   BIC = c(BIC(fit_1), BIC(fit_2), BIC(auto_fit))
12 )
13 print(model_comparison)

```

Listing 8: Model Selection and Diagnostics

```

1 checkresiduals(fit_1)
2 checkresiduals(fit_2)
3 checkresiduals(auto_fit)

```

Listing 9: 14-day Forecast

```
1 h <- 14
2 forecast_result <- forecast(auto_fit, h = h)
```

Listing 10: Plot of 14-day Forecast

```
1 autoplot(forecast_result) +
2   labs(
3     title = "14-Day Forecast of Average Uber Trip Duration",
4     x = "Day",
5     y = "Avg Trip Duration (minutes)"
6   ) +
7   theme_minimal()
```

Listing 11: 14-day Forecast Residuals

```
1 checkresiduals(forecast_result)
```

Listing 12: Table of Forecasted Values

```
1 start_date <- as.Date("2025-01-01")
2 forecast_dates <- seq.Date(from = start_date, by = "day", length.out = h)
3 forecast_summary <- as.data.frame(forecast_result)
4 formatted_forecast <- cbind(Date = forecast_dates, forecast_summary)
5
6 formatted_forecast <- within(formatted_forecast, {
7   'Point Forecast' <- round('Point Forecast', 2)
8   'Lo 80' <- round('Lo 80', 2)
9   'Hi 80' <- round('Hi 80', 2)
10  'Lo 95' <- round('Lo 95', 2)
11  'Hi 95' <- round('Hi 95', 2)
12 })
13
14 print(formatted_forecast)
```

C PySpark Preprocessing Code

The following PySpark script was used to preprocess over 200 million records from the NYC FHV dataset and compute the average daily trip durations.

First, the parquet file is loaded (Listing 13).

Listing 13: Loading FHV Data from Parquet Files

```
1 !pip install pyspark
2 from pyspark.sql import SparkSession
3 from pyspark.sql.functions import col, when, to_date, avg, round
4
5 spark = SparkSession.builder.master("local[*]").appName("FHV-ts-project").getOrCreate()
6
7 data_directory = "/content/drive/MyDrive/datasets/nyc-taxi/"
8 df = spark.read.format("parquet").option("recursiveFileLookup", "true").load(data_directory)
```

Derived columns are added for the duration of the trip in minutes and the average speed (Listing 14).

Listing 14: Deriving Trip Duration and Average Speed Columns

```
1 df = df.withColumn("trip_duration_min", col("trip_time") / 60)
2 df = df.withColumn("avg_speed_mph", (col("trip_miles") / (col("trip_time") / 3600)))
3
```

The data is then cleaned (Listing 15):

- Trips with zero or negative trip durations are removed.

- Outliers are removed:
 - Trips longer than 2 hours and shorter than 1 minute.
 - Trips longer than 100 miles and less than 0.1 miles.
 - Average speed of greater than 80 mph and less than 1 mph.

Listing 15: Data cleaning code

```

1 df_clean = df.filter(
2     (col("trip_duration_min") > 0) &
3     (col("trip_miles") > 0) &
4     (col("avg_speed_mph") > 0)
5 )
6
7
8 df_clean = df_clean.filter(
9     (col("trip_duration_min") <= 120) &
10    (col("trip_duration_min") >= 1) &
11    (col("trip_miles") <= 100) &
12    (col("trip_miles") >= 0.1) &
13    (col("avg_speed_mph") <= 80) &
14    (col("avg_speed_mph") >= 1)
15 )

```

The outliers are then removed (Listing 16).

Listing 16: Removing outliers code

```

1 columns_to_check = ["trip_duration_min", "trip_miles", "avg_speed_mph"]
2 iqr_dict = {}
3
4
5 for col_name in columns_to_check:
6     q1, q3 = df_clean.approxQuantile(col_name, [0.25, 0.75], 0.01)
7     iqr = q3 - q1
8     iqr_dict[col_name] = (q1, q3, iqr)
9     print(f"IQR_{col_name}: {iqr_dict[col_name][2]}")
10    print(f"Q1_{col_name}: {iqr_dict[col_name][0]}")
11    print(f"Q3_{col_name}: {iqr_dict[col_name][1]}")
12
13 iqr_duration = iqr_dict["trip_duration_min"][2]
14 iqr_miles = iqr_dict["trip_miles"][2]
15 iqr_speed = iqr_dict["avg_speed_mph"][2]
16
17 iqr_factor = 1.5
18
19 iqr_duration_15 = iqr_factor * iqr_duration
20 iqr_miles_15 = iqr_factor * iqr_miles
21 iqr_speed_15 = iqr_factor * iqr_speed
22
23 duration_upper = iqr_dict["trip_duration_min"][1] + iqr_duration_15
24 miles_upper = iqr_dict["trip_miles"][1] + iqr_miles_15
25 speed_upper = iqr_dict["avg_speed_mph"][1] + iqr_speed_15
26
27 df_final = df_clean.filter(
28     (col("trip_duration_min") <= duration_upper) &
29     (col("trip_miles") <= miles_upper) &
30     (col("avg_speed_mph") <= speed_upper)
31 )
32

```

Keep only Uber trips (Listing 17).

Listing 17: Filter to Uber trips code

```

1 df_uber = df.filter(df["hvfhs_license_num"] == "HV0003")

```

Then the average trip duration per day is calculated (Listing 18).

Listing 18: Calculate average daily duration code

```

1 df_uber = df_uber.withColumn("pickup_date", to_date(col("pickup_datetime")))
2
3 df_daily_avg = df_uber.groupBy("pickup_date") \
4     .agg(avg("trip_duration_min").alias("avg_duration_min")) \
5     .orderBy("pickup_date")
6
7 df_daily_avg = df_daily_avg.withColumn("avg_duration_min", round("avg_duration_min"))
8

```


Finally, the cleaned time series dataset is saved in a CSV file (Listing 19). The resulting CSV file contains two columns, the pickup date, and average trip duration for that day.

Listing 19: Saving time series data code

```
1 final_path = "/content/drive/MyDrive/datasets/fhvhv_tripdata_2024_uber_timeseries.csv"
2
3
4 df_daily_avg.select("pickup_date", "avg_duration_min") \
5     .coalesce(1) \
6     .write.option("header", True) \
7     .csv(final_path)
```

References

- [1] Jiaying Ma, Wei Zhou, Fan Wang, et al. Congestion-aware electric vehicle scheduling for ride-hailing fleets. arXiv preprint arXiv:2412.09978, 2024.
- [2] New York City Taxi and Limousine Commission. Tlc trip record data. <https://www.nyc.gov/site/tlc/about/tlc-trip-record-data.page>, 2024. Accessed July 2025.
- [3] Tandon School of Engineering. Nyc’s ride-hailing fee failed to ease manhattan traffic, new nyu tandon study reveals. NYU Tandon News, 2024. Accessed July 2025.