

# HW5

Farooq Mahmud

## Problem 1

```
library(forecast)
```

```
Registered S3 method overwritten by 'quantmod':  
  method      from  
  as.zoo.data.frame zoo
```

```
library(TSA)
```

```
Registered S3 methods overwritten by 'TSA':  
  method      from  
  fitted.Arima forecast  
  plot.Arima   forecast
```

```
Attaching package: 'TSA'
```

```
The following objects are masked from 'package:stats':
```

```
  acf, arima
```

```
The following object is masked from 'package:utils':
```

```
  tar
```

```
library(ggplot2)
data(robot)

model_ar1 <- Arima(robot, order = c(1, 0, 0))
model_arima011 <- Arima(robot, order = c(0, 1, 1))


model_comparison <- data.frame(
  Model = c("AR(1)", "ARIMA(0,1,1)"),
  AIC = c(AIC(model_ar1), AIC(model_arima011)),
  BIC = c(BIC(model_ar1), BIC(model_arima011))
)

print(model_comparison)
```

	Model	AIC	BIC
1	AR(1)	-2945.078	-2933.735
2	ARIMA(0,1,1)	-2957.901	-2950.346

```
checkresiduals(model_ar1)
```

Residuals from ARIMA(1,0,0) with non-zero mean




Ljung-Box test

```
data: Residuals from ARIMA(1,0,0) with non-zero mean  
Q* = 42.118, df = 9, p-value = 3.127e-06
```

```
Model df: 1. Total lags used: 10
```

```
checkresiduals(model_arima011)
```



Ljung-Box test

```
data: Residuals from ARIMA(0,1,1)  
Q* = 14.796, df = 9, p-value = 0.0967
```

```
Model df: 1. Total lags used: 10
```


Based on lower AIC/BIC values and better residual diagnostics (especially the Ljung-Box test), ARIMA(0,1,1) is the preferred model.

## Problem 2

```
data(deere3)
model_ar1 <- Arima(deere3, order = c(1, 0, 0))

forecast_ar1 <- forecast::forecast(model_ar1, h = 10)

autoplot(forecast_ar1) +
  labs(
    title = "Forecast from AR(1) Model for deere3"
  ) +
  theme_minimal()
```



```
forecast_table <- as.data.frame(forecast_ar1)

print(forecast_table)
```

	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
58	-335.145928	-2211.910	1541.618	-3205.409	2535.117
59	-117.120772	-2237.282	2003.041	-3359.628	3125.386
60	-2.538388	-2185.148	2180.071	-3340.551	3335.474

61	57.679997	-2141.865	2257.225	-3306.233	3421.593
62	89.327566	-2114.872	2293.527	-3281.705	3460.360
63	105.959839	-2099.523	2311.443	-3267.036	3478.955
64	114.700873	-2091.137	2320.539	-3258.837	3488.239
65	119.294695	-2086.641	2325.230	-3254.393	3492.982
66	121.708962	-2084.254	2327.672	-3252.020	3495.438
67	122.977772	-2082.992	2328.948	-3250.762	3496.718

The forecast quickly reverts toward zero (the series mean), consistent with the nature of AR(1) models. The wide prediction intervals indicate substantial uncertainty, which is expected due to the variability observed in the original series.

### Problem 3a



```
library(tseries)

gdp <- read.csv("CAN_GDP.csv")
int <- read.csv("CAN_INT.csv")
cpi <- read.csv("CAN_CPI.csv")
pro <- read.csv("CAN_PRO.csv")



ts_gdp <- ts(gdp[, 7])
ts_int <- ts(int[, 7])
ts_cpi <- ts(cpi[, 7])
ts_pro <- ts(pro[, 7])

plot_acf_pacf <- function(ts_data, title) {
  par(mfrow = c(1, 2))
  acf(ts_data, main = paste("ACF of", title))
  pacf(ts_data, main = paste("PACF of", title))
  par(mfrow = c(1, 1))
}

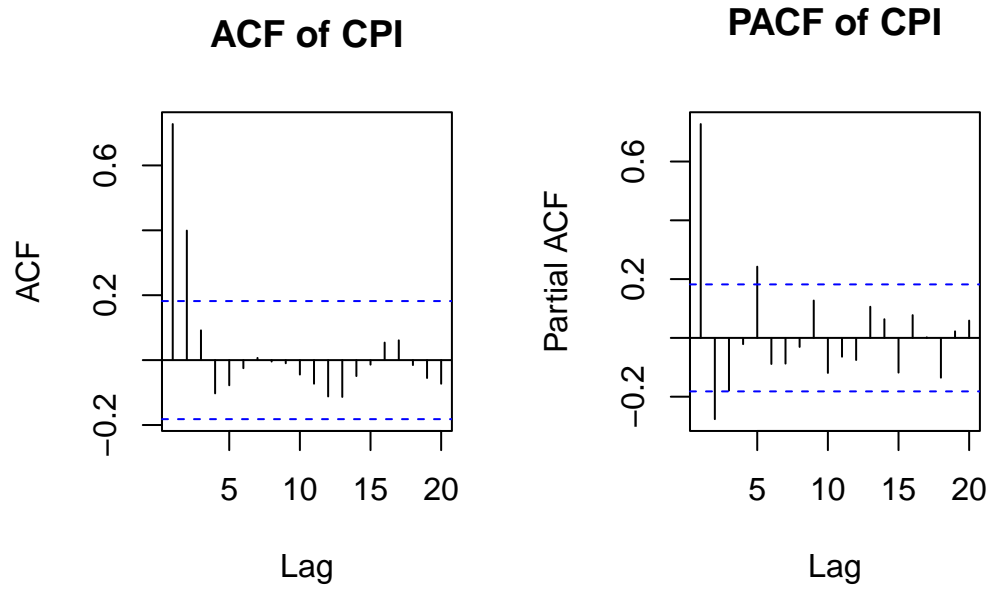
plot_acf_pacf(ts_gdp, "GDP")
```



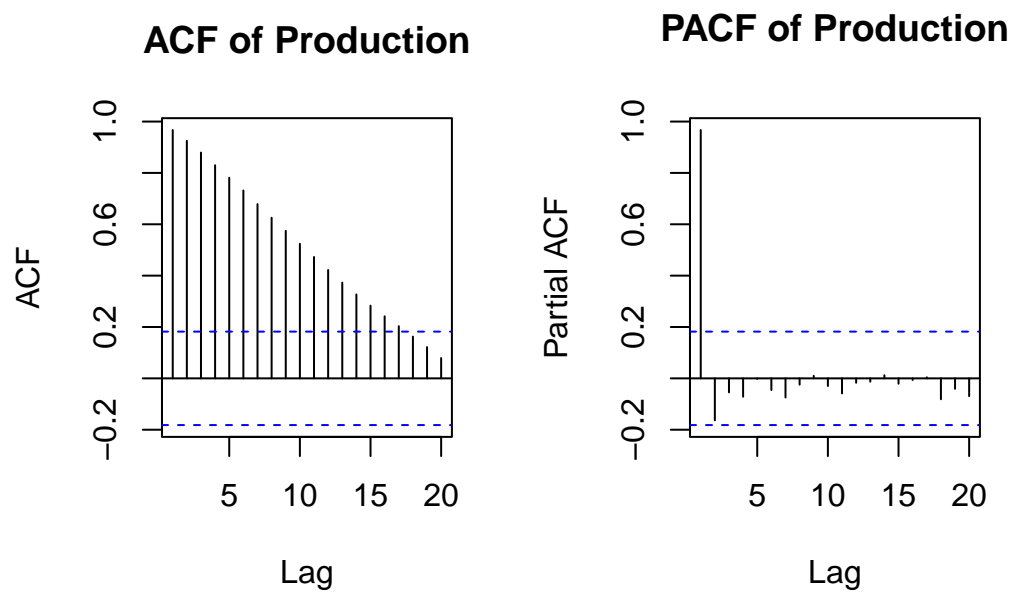
```
plot_acf_pacf(ts_int, "Interest Rate")
```



```
plot_acf_pacf(ts_cpi, "CPI")
```



```
plot_acf_pacf(ts_pro, "Production")
```



## Stationarity Assessments


- GDP
  - The ACF shows a slow, exponential decay, typical of a non-stationary time series.
  - The PACF cuts off sharply after lag 1, typical of an AR(1) process after differencing, suggesting the underlying data could be modeled as an ARIMA(1,1,0) process.
- Interest Rate
  - Same assessment as GDP.
- CPI
  - ACF shows a large spike at lag 1, followed by quickly diminishing values which is a sign of stationarity.
  - The PACF cuts off sharply after lag 1, typical of an AR(1) process but differencing may not be required.
- Production
  - Same assessment as GDP and Interest Rate.

## Problem 3b


```
log_diff <- function(series) {  
  diff(log(series))  
}  
  
ts_gdp_diff <- log_diff(ts_gdp)  
ts_int_diff <- log_diff(ts_int)  
ts_pro_diff <- log_diff(ts_pro)  
  
plot_acf_pacf(ts_gdp_diff, "Log-Diff GDP")
```



**ACF of Log-Diff GDP**





**PACF of Log-Diff GDP**




```
plot_acf_pacf(ts_int_diff, "Log-Diff Interest Rate")
```

**ACF of Log-Diff Interest Rate**      **PACF of Log-Diff Interest Rate**



```
plot_acf_pacf(ts_pro_diff, "Log-Diff Production")
```

## ACF of Log-Diff Production      PACF of Log-Diff Productic



### Stationarity Assessments

- GDP
  - Only Lag 1 ACF is significant.
  - In the PACF, only the first lag is significant.
  - The transformed series appears stationary, possibly an AR(1) process.
- Interest Rate
  - ACF and PACF plots show most autocorrelations are within the confidence bounds.
  - The transformed series resembles white noise more closely.
- Production
  - Transformed series appears stationary for the same reasons as GDP.

### Problem 3c

```
library(vars)
```

Loading required package: MASS

Loading required package: strucchange

Loading required package: zoo

Attaching package: 'zoo'

The following objects are masked from 'package:base':

as.Date, as.Date.numeric

Loading required package: sandwich

Loading required package: urca

Loading required package: lmtest

```
clean_ts <- function(ts_data) {  
  vec <- as.numeric(ts_data)  
  vec[!is.finite(vec)] <- NA  
  vec <- na.omit(vec)  
  ts(vec, start = start(ts_data), frequency = frequency(ts_data))  
}  
  
ts_cpi_diff <- log_diff(ts_cpi)  
ts_cpi_diff <- clean_ts(ts_cpi_diff) # Differenced CPI has INF, -INF. This gets rid of them.  
  
combined_diff <- cbind(ts_gdp_diff, ts_int_diff, ts_cpi_diff, ts_pro_diff)  
combined_diff <- na.omit(combined_diff) # Remove NA's.  
  
VARselect(combined_diff)
```

```

$selection
AIC(n)  HQ(n)  SC(n) FPE(n)
      1      1      1      1

$criteria
      1      2      3      4      5
AIC(n) -2.567514e+01 -2.557521e+01 -2.550112e+01 -2.553925e+01 -2.554834e+01
HQ(n)   -2.546550e+01 -2.519786e+01 -2.495606e+01 -2.482647e+01 -2.466786e+01
SC(n)   -2.515729e+01 -2.464308e+01 -2.415472e+01 -2.377857e+01 -2.337339e+01
FPE(n)   7.072460e-12  7.828027e-12  8.462532e-12  8.204730e-12  8.225907e-12
      6      7      8      9     10
AIC(n) -2.541890e+01 -2.528148e+01 -2.524889e+01 -2.512759e+01 -2.500693e+01
HQ(n)   -2.437071e+01 -2.406558e+01 -2.386527e+01 -2.357627e+01 -2.328789e+01
SC(n)   -2.282967e+01 -2.227798e+01 -2.183111e+01 -2.129553e+01 -2.076060e+01
FPE(n)   9.526864e-12  1.120079e-11  1.196037e-11  1.410061e-11  1.681873e-11

```

Based on the above output, VAR(1) is the most appropriate.

### Problem 3d

```

var_model <- VAR(combined_diff, p = 1)
summary(var_model)

```

VAR Estimation Results:

=====

Endogenous variables: ts\_gdp\_diff, ts\_int\_diff, ts\_cpi\_diff, ts\_pro\_diff

Deterministic variables: const

Sample size: 110

Log Likelihood: 814.277

Roots of the characteristic polynomial:

0.5339 0.3401 0.1923 0.1062

Call:

VAR(y = combined\_diff, p = 1)

Estimation results for equation ts\_gdp\_diff:

=====

ts\_gdp\_diff = ts\_gdp\_diff.l1 + ts\_int\_diff.l1 + ts\_cpi\_diff.l1 + ts\_pro\_diff.l1 + const

	Estimate	Std. Error	t value	Pr(> t )
ts_gdp_diff.l1	0.3268818	0.1250852	2.613	0.0103 *
ts_int_diff.l1	0.0137291	0.0051736	2.654	0.0092 **
ts_cpi_diff.l1	-0.0020123	0.0009251	-2.175	0.0319 *
ts_pro_diff.l1	0.0635530	0.0425058	1.495	0.1379
const	0.0040070	0.0008061	4.971	2.6e-06 ***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.004894 on 105 degrees of freedom

Multiple R-Squared: 0.3541, Adjusted R-squared: 0.3295

F-statistic: 14.39 on 4 and 105 DF, p-value: 2.114e-09

Estimation results for equation ts\_int\_diff:

=====

ts\_int\_diff = ts\_gdp\_diff.l1 + ts\_int\_diff.l1 + ts\_cpi\_diff.l1 + ts\_pro\_diff.l1 + const

	Estimate	Std. Error	t value	Pr(> t )
ts_gdp_diff.l1	6.740654	2.309444	2.919	0.00430 **
ts_int_diff.l1	0.001691	0.095520	0.018	0.98591
ts_cpi_diff.l1	0.002474	0.017080	0.145	0.88513
ts_pro_diff.l1	-2.085685	0.784782	-2.658	0.00910 **
const	-0.046362	0.014882	-3.115	0.00237 **

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.09035 on 105 degrees of freedom

Multiple R-Squared: 0.0804, Adjusted R-squared: 0.04537

F-statistic: 2.295 on 4 and 105 DF, p-value: 0.06408

Estimation results for equation ts\_cpi\_diff:

=====

ts\_cpi\_diff = ts\_gdp\_diff.l1 + ts\_int\_diff.l1 + ts\_cpi\_diff.l1 + ts\_pro\_diff.l1 + const

	Estimate	Std. Error	t value	Pr(> t )
ts_gdp_diff.l1	0.16196	12.71032	0.013	0.98986
ts_int_diff.l1	0.72471	0.52571	1.379	0.17097
ts_cpi_diff.l1	-0.28988	0.09400	-3.084	0.00261 **
ts_pro_diff.l1	-0.43724	4.31915	-0.101	0.91956

```

const          -0.01479    0.08191  -0.181  0.85702
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

Residual standard error: 0.4972 on 105 degrees of freedom  
Multiple R-Squared: 0.09562, Adjusted R-squared: 0.06117  
F-statistic: 2.775 on 4 and 105 DF, p-value: 0.03076

Estimation results for equation ts\_pro\_diff:

```

=====
ts_pro_diff = ts_gdp_diff.l1 + ts_int_diff.l1 + ts_cpi_diff.l1 + ts_pro_diff.l1 + const

```

	Estimate	Std. Error	t value	Pr(> t )
ts_gdp_diff.l1	0.912259	0.379862	2.402	0.0181 *
ts_int_diff.l1	0.025752	0.015711	1.639	0.1042
ts_cpi_diff.l1	-0.003897	0.002809	-1.387	0.1684
ts_pro_diff.l1	0.241205	0.129083	1.869	0.0645 .
const	-0.002635	0.002448	-1.076	0.2842

```

---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

Residual standard error: 0.01486 on 105 degrees of freedom  
Multiple R-Squared: 0.3271, Adjusted R-squared: 0.3015  
F-statistic: 12.76 on 4 and 105 DF, p-value: 1.683e-08

Covariance matrix of residuals:

	ts_gdp_diff	ts_int_diff	ts_cpi_diff	ts_pro_diff
ts_gdp_diff	2.395e-05	7.631e-05	-0.0003236	4.887e-05
ts_int_diff	7.631e-05	8.163e-03	0.0033333	2.837e-04
ts_cpi_diff	-3.236e-04	3.333e-03	0.2472568	-1.531e-03
ts_pro_diff	4.887e-05	2.837e-04	-0.0015309	2.208e-04

Correlation matrix of residuals:

	ts_gdp_diff	ts_int_diff	ts_cpi_diff	ts_pro_diff
ts_gdp_diff	1.0000	0.17259	-0.13298	0.6719
ts_int_diff	0.1726	1.00000	0.07419	0.2113
ts_cpi_diff	-0.1330	0.07419	1.00000	-0.2072
ts_pro_diff	0.6719	0.21130	-0.20717	1.0000

### Problem 3e

```
fit_ar_model <- function(ts_series, series_name) {  
  model <- arima(ts_series, order = c(1, 0, 0))  
  cat("\n==== AR(1) Model for", series_name, "====\n")  
  print(summary(model))  
  return(model)  
}  
  
ar_gdp <- fit_ar_model(ts_gdp_diff, "GDP")
```

==== AR(1) Model for GDP =====

Call:

```
arima(x = ts_series, order = c(1, 0, 0))
```

Coefficients:

	ar1	intercept
	0.4943	6e-03
s.e.	0.0804	9e-04

sigma<sup>2</sup> estimated as 2.589e-05: log likelihood = 443.98, aic = -883.95

Training set error measures:

Warning in trainingaccuracy(object, test, d, D): test elements must be within sample

	ME	RMSE	MAE	MPE	MAPE
Training set	NaN	NaN	NaN	NaN	NaN

```
ar_int <- fit_ar_model(ts_int_diff, "Interest Rate")
```

==== AR(1) Model for Interest Rate =====

Call:

```
arima(x = ts_series, order = c(1, 0, 0))
```

Coefficients:

	ar1	intercept
	0.0148	-0.0162
s.e.	0.0935	0.0090

sigma^2 estimated as 0.008952: log likelihood = 107.99, aic = -211.97

Training set error measures:

Warning in trainingaccuracy(object, test, d, D): test elements must be within sample

	ME	RMSE	MAE	MPE	MAPE
Training set	NaN	NaN	NaN	NaN	NaN

```
ar_cpi <- fit_ar_model(ts_cpi_diff, "CPI")
```

===== AR(1) Model for CPI =====

Call:

```
arima(x = ts_series, order = c(1, 0, 0))
```

Coefficients:

	ar1	intercept
	-0.2788	-0.0195
s.e.	0.0906	0.0363

sigma^2 estimated as 0.2382: log likelihood = -77.91, aic = 159.82

Training set error measures:

Warning in trainingaccuracy(object, test, d, D): test elements must be within sample

	ME	RMSE	MAE	MPE	MAPE
Training set	NaN	NaN	NaN	NaN	NaN



```
ar_pro <- fit_ar_model(ts_pro_diff, "Production")
```

===== AR(1) Model for Production =====

Call:

```
arima(x = ts_series, order = c(1, 0, 0))
```

Coefficients:

	ar1	intercept
	0.5100	0.0033
s.e.	0.0795	0.0028

sigma^2 estimated as 0.0002231: log likelihood = 320.11, aic = -636.23

Training set error measures:

Warning in trainingaccuracy(object, test, d, D): test elements must be within sample

	ME	RMSE	MAE	MPE	MAPE
Training set	NaN	NaN	NaN	NaN	NaN

### Problem 3f

```
calculate_ar_mse <- function(models, data_list) {  
  ar_mse <- sapply(seq_along(models), function(i) {  
    resid <- residuals(models[[i]])  
    mean(resid^2, na.rm = TRUE)  
  })  
  avg_mse <- mean(ar_mse)  
  total_params <- length(models) * 2 # Each AR(1): 1 lag + 1 intercept  
  list(avg_mse = avg_mse, total_params = total_params)  
}  
  
calculate_var_mse <- function(var_model) {  
  resid <- residuals(var_model)  
  
  # Remove rows with any NA/Inf across variables
```

```

resid_clean <- resid[apply(resid, 1, function(row) all(is.finite(row))), ]

# MSE per series
mse_per_series <- colMeans(resid_clean^2, na.rm = TRUE)
avg_mse <- mean(mse_per_series)

k <- ncol(resid) # Number of variables
p <- var_model$p # VAR order
total_params <- k^2 * p + k # per standard VAR(p) parameter count

list(avg_mse = avg_mse, total_params = total_params)
}

ar_models <- list(ar_gdp, ar_int, ar_cpi, ar_pro)
ar_results <- calculate_ar_mse(ar_models, list(ts_gdp_diff, ts_int_diff, ts_cpi_diff, ts_pro))
var_results <- calculate_var_mse(var_model)

table <- data.frame(
  Model = c("AR(1)", sprintf("VAR(%d)", var_model$p)),
  Avg_MSE = c(round(ar_results$avg_mse, 5), round(var_results$avg_mse, 5)),
  Num_Parameters = c(ar_results$total_params, var_results$total_params)
)

print(table)

```

	Model	Avg_MSE	Num_Parameters
1	AR(1)	0.06184	8
2	VAR(1)	0.06101	20

## Problem 5g

Although the VAR(1) model has a slightly lower average MSE than the AR(1) models, the difference is small. The VAR(1) model uses more parameters than the AR(1) models which could lead to concerns about model complexity and overfitting.

Therefore, AR(1) is recommended for parsimony unless inter-variable relationships are essential to capture.